

# Smarter Balanced Adaptive Item Selection Algorithm Design Report

V3. 5/9/2014

**Jon Cohen and Larry Albright**  
**American Institutes for Research**

Produced for Smarter Balanced by the American  
Institutes for Research

This work is licensed under a [Creative Commons  
Attribution-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nd/4.0/).



## TABLE OF CONTENTS

1. INTRODUCTION, BACKGROUND, AND DEFINITIONS.....	1
1.1 Blueprint .....	2
1.2 Content Value .....	3
1.2.1 Content Value for Single Items.....	3
1.2.2 Content Value for Sets of Items.....	4
1.3 Information Value.....	6
1.3.1 Individual Information Value .....	6
1.3.2 Binary Items.....	6
1.3.3 Polytomous Items.....	6
1.3.4 Item Group Information Value .....	9
2. ENTRY AND INITIALIZATION.....	9
2.1 Item Pool Selection .....	9
2.2 Adjust Segment Length.....	9
2.3 Initialization of Starting Theta Estimates .....	10
2.4 Insertion of Embedded Field-Test Items.....	10
3. ITEM SELECTION .....	11
3.1 Trimming the Custom Item Pool .....	12
3.2 Recycling Algorithm.....	13
3.3 Adaptive Item Selection.....	13
3.4 Selection of the Initial Item.....	14
3.5 Exposure Control.....	12
4. TERMINATION.....	14
A1. DEFINITIONS OF USER-SETTABLE PARAMETERS.....	15
A2. API .....	16
A3. SUPPORTING DATA STRUCTURES .....	16

## SMARTER BALANCED ADAPTIVE ITEM SELECTION ALGORITHM

### 1. INTRODUCTION, BACKGROUND, AND DEFINITIONS

This document describes the Smarter Balanced adaptive item selection algorithm. The item selection algorithm is designed to cover a standards-based blueprint, which may include content, cognitive complexity, and item type constraints. The item selection algorithm will also include:

- the ability to customize an item pool based on access constraints and screen items that have been previously viewed or may not be accessible for a given individual;
- a mechanism for inserting embedded field-test items; and
- a mechanism for delivering “segmented” tests in which separate parts of the test are administered in a fixed order.

This document describes the algorithm and the design for its implementation for the Smarter Balanced Test Delivery System. The implementation builds extensively on the algorithm implemented in AIR’s Test Delivery System. The implementation described is released under the Creative Commons Attribution Only, No Derivatives license.

The general approach described here is based on a highly parameterized multiple-objective utility function. The objective function includes:

- a measure of content match to the blueprint;
- a measure of overall test information; and
- measures of test information for each reporting category on the test.

We define an objective function that measures an item’s contribution to each of these objectives, weighting them to achieve the desired balance among them. Equation 1 sketches this objective function for a single item.

$$f_{ijt} = w_2 \frac{1}{\sum_{r=1}^R d_{rj}} \sum_{r=1}^R s_{rit} p_r d_{rj} + w_1 \sum_{k=1}^K q_k h_k(v_{kijt}, V_{kit}, t_k) + w_0 h_0(u_{ijt}, U_{it}, t_0) \quad (1)$$

where the terms  $w$  represent user-supplied weights that assign relative importance to meeting each of the objectives,  $d_{rj}$  indicates whether item  $j$  has the blueprint-specified feature  $r$ , and  $p_r$  is the user-supplied priority weight for feature  $r$ . The term  $s_{rit}$  is an adaptive control parameter that is described below. In general,  $s_{rit}$  increases for features that have not met their designated minimum as the end of the test approaches.

The remainder of the terms represents an item’s contribution to measurement precision:

- $v_{kijt}$  is the value of item  $j$  toward reducing the measurement error for reporting category  $k$  for examinee  $i$  at selection  $t$ ; and

- $u_{ijt}$  is the value of item  $j$  in terms of reducing the overall measurement error for examinee  $i$  at selection  $t$ .

The terms  $U_{it}$  and  $V_{kit}$  represent the total information overall and on reporting category  $k$ , respectively.

The term  $q_k$  is a user-supplied priority weight associated with the precision of the score estimate for reporting category  $k$ . The terms  $t$  represent precision targets for the overall score ( $t_0$ ) and each score reporting category score. The functions  $h(\cdot)$  are given by:

$$h_0(u_{ijt}, U_{it}, t_0) = \begin{cases} au_{ijt} & \text{if } U_{it} < t_0 \\ bu_{ijt} & \text{otherwise} \end{cases}$$

$$h_{1k}(v_{kijt}, V_{kit}, t_k) = \begin{cases} c_k v_{kijt} & \text{if } V_{kit} < t_k \\ d_k v_{kijt} & \text{otherwise} \end{cases}$$

Items can be selected to maximize the value of this function. This objective function can be manipulated to produce a pure, standards-free adaptive algorithm by setting  $w_2$  to zero or a completely blueprint-driven test by setting  $w_1 = w_0 = 0$ . Adjusting the weights to optimize performance for a given item pool will enable users to maximize information subject to the constraint that the blueprint is virtually always met.

We note that the computations of the content values and information values generate values on very different scales and that the scale of the content value varies as the test progresses. Therefore, we normalize both the information and content values before computing the value of

Equation 1. This normalization is given by  $x = \begin{cases} 1 & \text{if } \min = \max \\ \frac{v - \min}{\max - \min} & \text{otherwise} \end{cases}$ , where  $\min$  and  $\max$  represent the minimum and maximum, respectively, of the metric computed over the current set of items or item groups.

The remainder of this section describes the overall program flow, the form of the blueprint, and the various value calculations employed in the objective function. Subsequent sections describe the details of the selection algorithm.

## 1.1 Blueprint

Each test will be described by a single blueprint for each segment of the test and will identify the order in which the segments appear. The blueprint will include:

- an indicator of whether the test is adaptive or fixed form;
- termination conditions for the segment, which are described in a subsequent section;
- a set of nested content constraints, each of which is expressed as:
  - the minimum number of items to be administered within the content category;
  - the maximum number of items to be administered within the content category;

- an indication of whether the maximum should be deterministically enforced (a “strict” maximum);
- a priority weight for the content category  $p_r$ ;
- an explicit indicator as to whether this content category is a reporting category; and
- an explicit precision-priority weight ( $q_k$ ) for each group identified as a reporting category.
- a set of non-nested content constraints, which are represented as:
  - a name for the collection of items meeting the constraint;
  - the minimum number of items to be administered from this group of items;
  - the maximum number of items to be administered from this group of items;
  - an indication of whether the maximum should be deterministically enforced (a “strict” maximum);
  - a priority weight for the group of items  $p_r$ ;
  - an explicit indicator as to whether this named group will make up a reporting category; and
  - an explicit precision-priority weight ( $q_k$ ) for each group identified as a reporting category.
- The priority weights,  $p_r$  on the blueprint, can be used to express values in the blueprint match. Large weights on reporting categories paired with low (or zero) weights on the content categories below them may allow more flexibility to maximize information in a content category covering fewer fine-grained targets, while the reverse would mitigate toward more reliable coverage of finer-grained categories, with less content flexibility within reporting categories.

An example of a blueprint specification appears in Appendix 1.

Each segment of a test will have a separate blueprint.

## 1.2 Content Value

Each item or item group will be characterized by its contribution to meeting the blueprint, given the items that have already been administered at any point. The contribution is based on the presence or absence of features specified in the blueprint and denoted by the term  $d$  in Equation 1. This section describes the computation of the content value.

### 1.2.1 Content Value for Single Items

For each constraint appearing in the blueprint ( $r$ ), an item  $i$  either does or does not have the characteristic described by the constraint. For example, a constraint might require a minimum of four and a maximum of six algebra items. An item measuring algebra has the described

characteristic, and an item measuring geometry but algebra does not. To capture this constraint, we define the following:

- $d_i$  is a feature vector in which the elements are  $d_{ir}$ , summarizing item  $i$ 's contribution to meeting the blueprint. This feature vector includes content categories such as claims and targets as well as other features of the blueprint, such as Depth of Knowledge and item type.
- $S_{it}$  is a diagonal matrix, the diagonal elements of which are the adaptive control parameters  $s_{rit}$ .
- $p$  is the vector containing the user-supplied priority weights  $p_r$ .

The scalar content value for an item is given by  $C_{ijt} = d_i' S_{it} p$ .

Letting  $z_{rit}$  represent the number of items with feature  $r$  administered to student  $i$  by iteration  $t$ , the value of the adaptive control parameters is:

$$s_{rit} = \begin{cases} m_{it} \left( 2 - \frac{z_{rit}}{Min_r} \right) & \text{if } z_r < Min_r \\ 1 - \frac{z_{rit} - Min_r}{Max_r - Min_r} & \text{if } Min_r < z_{rit} < Max_r \\ (Max_r - z_{rit}) - 1 & \text{if } Max_r \leq z_{rit} \end{cases}$$

The blueprint defines the minimum ( $Min_r$ ) and maximum ( $Max_r$ ) number of items to be administered with each characteristic ( $r$ ).

The term  $m_{it} = \frac{T}{T-t}$  where  $T$  is the total test length. This has the effect of increasing the algorithm's preference for items that have not yet met their minimums as the end of the test nears and the opportunities to meet the minimum diminish.

This increases the likelihood of selecting items for content that has not met its minimum as the opportunities to do so are used up. The value  $s$  is highest for items with content that has not met its minimum, declines for items representing content for which the minimum number of items has been reached but the maximum has not, and turns negative for items representing content that has met the maximum.

## 1.2.2 Content Value for Sets of Items

Calculation of the content value of sets of items is complicated by two factors:

1. The desire to allow more items to be developed for each set and to have the most advantageous set of items administered

2. The design objective of characterizing the information contribution of a set of items as the expected information over the working theta distribution for the examinee

The former objective is believed to enhance the ability to satisfy highly constrained blueprints while still adapting to obtain good measurement for a broad range of students. The latter arises from the recognition that ELA tests will select one set of items at a time, without an opportunity to adapt once the passage has been selected.

The general approach involves successive selection of the highest content value item in the set until the indicated number of items in the set have been selected. Because the content value of an item changes with each selection, a temporary copy of the already-administered content vector for the examinee is updated with each selection such that subsequent selections reflect the items selected in previous iterations.

Exhibit 1 presents a flowchart for this calculation. Readers will note the check to determine whether  $w_0 > 0$  or  $w_1 > 0$ . These weights, defined with Equation 1, identify the user-supplied importance of information optimization relative to blueprint optimization. In cases such as independent field tests, this weight may be set to zero, as it may not be desirable to make item administration dependent on match to student performance. In more typical adaptive cases where item statistics will not be recalculated, favoring more informative items is generally better. The final measure of content value for the set of selected set of items is divided by the number of items selected to avoid a bias toward selection of sets with more items.

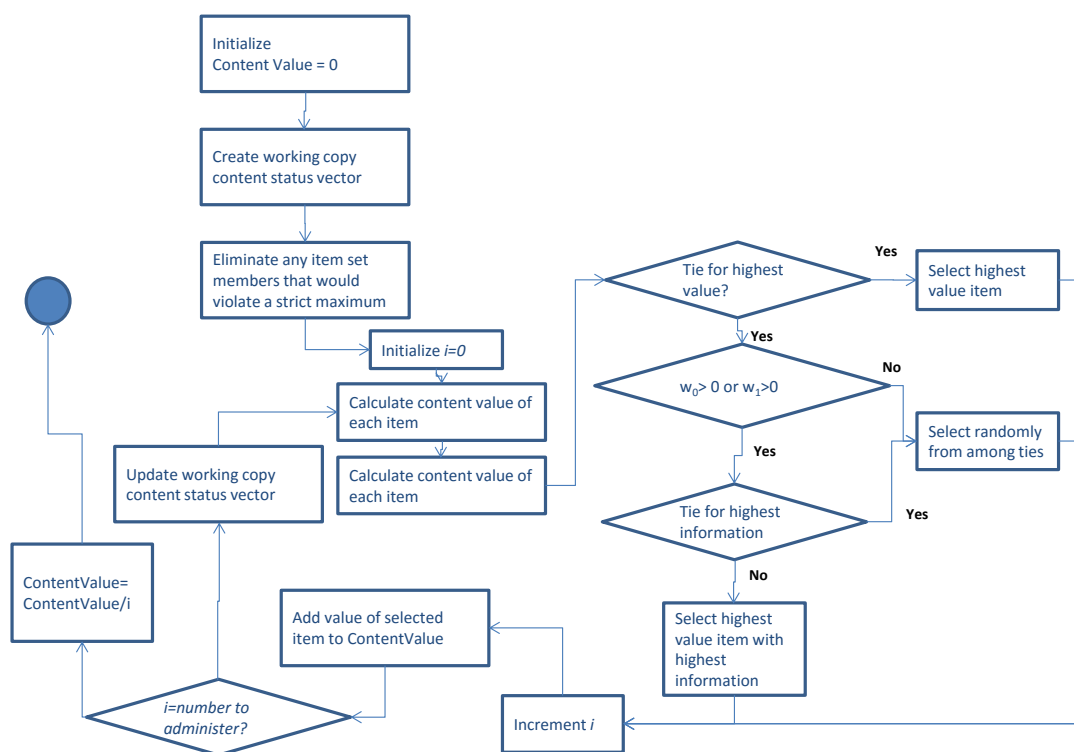


Exhibit 1. Content Value Calculation for Item Sets

## 1.3 Information Value

Each item or item group also has value in terms of maximizing information, both overall and on reporting categories.

### 1.3.1 Individual Information Value

The information value associated with an item will be an approximation of information. The system will be designed to use generalized IRT models; however, it will treat all items as though they offer equal measurement precision. This is the assumption made by the Rasch model, but in more general models, items known to offer better measurement are given preference by many algorithms. Subsequent algorithms are then required to control the exposure of the items that measure best. Ignoring the differences in slopes serves to eliminate this bias and help equalize exposure.

### 1.3.2 Binary Items

The approximate information value of a binary item will be characterized as  $I_j(\theta) = p_j(\theta)(1 - p_j(\theta))$ , where the slope parameters are artificially replaced with a constant.

### 1.3.3 Polytomous Items

In terms of information, the best polytomous item in the pool is the one that maximizes the expected information,  $I_j(\theta)$ . Formally,  $I_j(\theta) > I_k(\theta)$  for all items  $k \neq j$ . The true value  $\theta$ , however, remains unknown and is accessed only through an estimate,  $\hat{\theta} \sim N(\bar{\theta}, \sigma_{\theta})$ . By definition of an expectation, the expected information  $I_j(\theta) = \int I_j(t)f(t|\bar{\theta}, \sigma_{\theta})dt$ .

The intuition behind this result is illustrated in Exhibit 2. In Exhibit 2, each panel graphs the distribution of the estimate of  $\theta$  for an examinee. The top panel assumes a polytomous item in which one step threshold (A1) matches the mean of the  $\theta$  estimate distribution. In the bottom panel, neither step threshold matches the mean of the  $\theta$  estimate distribution. The shaded area in each panel indicates the region in which the hypothetical item depicted in the panel provides more information. We see that approximately 2/3 of the probability density function is shaded in the lower panel, while the item depicted in the upper panel dominates in only about 1/3 of the cases. In this example, the item depicted in the lower panel has a much greater probability of maximizing the information from the item, despite the fact that the item in the upper panel has a threshold exactly matching the mean of the estimate distribution and the item in the lower panel does not.




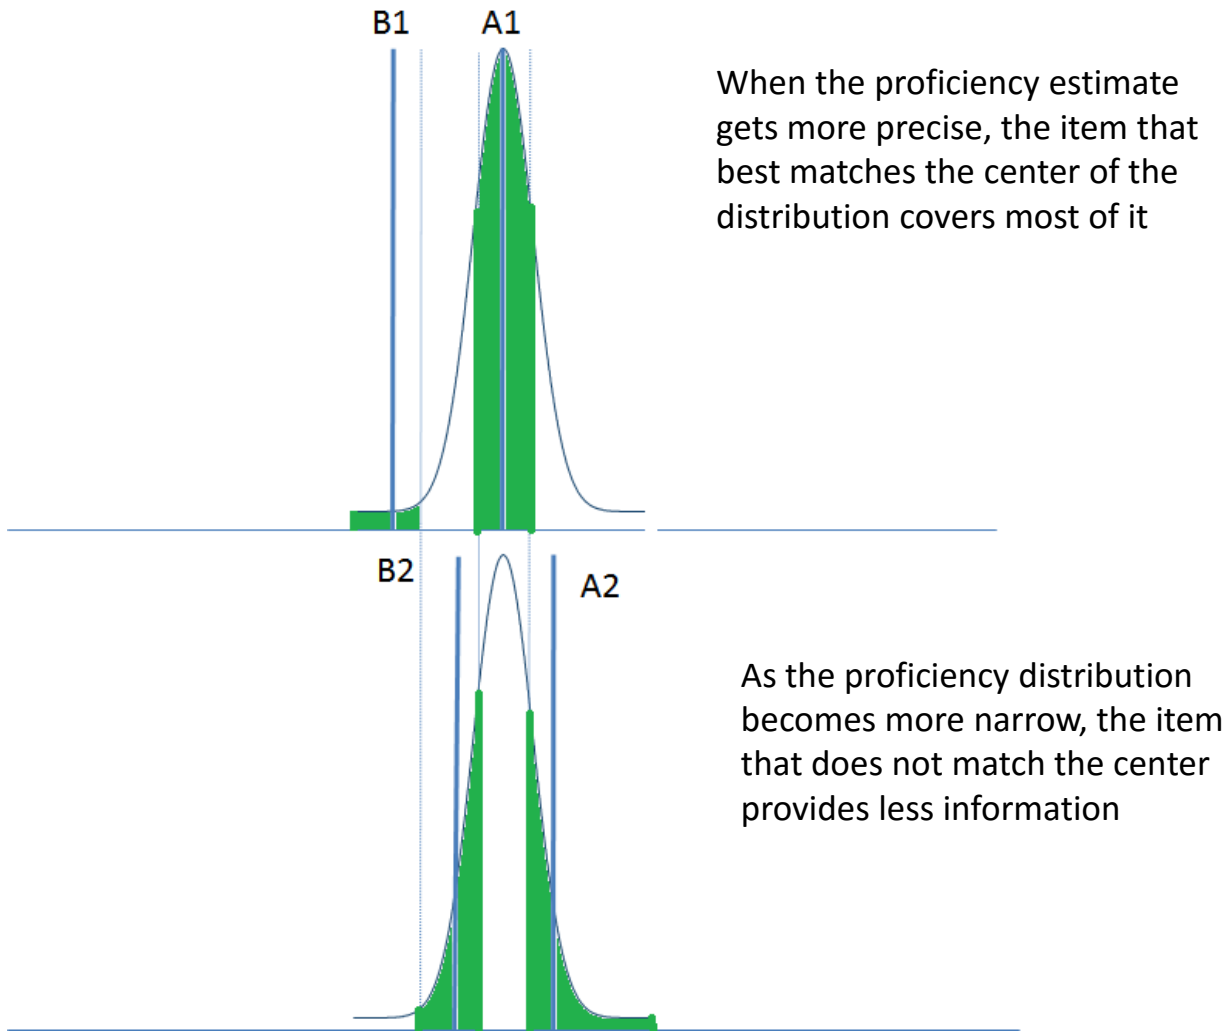


Exhibit 2. Two example items, with the shaded region showing the probability that the item maximizes information for the examinee depicted.

Exhibit 3 shows what happens to information as the estimate of this student's proficiency becomes more precise (later in the test). In this case, the item depicted in the top panel maximizes information about 65-70 percent of the time, compared to about 30 to 35 percent for the item depicted in the lower panel. These are the same items depicted in the Exhibit 2, but in this case we are considering information for a student with a more precise current proficiency estimate.



**Exhibit 3: Two example items, with the shaded region showing the probability that the item maximizes information for the examinee depicted.**

The approximate information value of polytomous items will be characterized as the expected information, specifically  $E[I_j(\theta)|m_i, s_i] = \int \sum_{k=1}^K I_{jk}(t) p_j(k|t) \phi(t; m_i, s_i) dt$ , where  $I_{jk}(t)$  represents the information at  $t$  of response  $k$  to item  $j$ ,  $p_j(k|t)$  is the probability of response  $k$  to item  $j$  (artificially holding slope constant), given proficiency  $t$ ,  $\phi(\cdot)$  represents the normal probability density function, and  $m_i$  and  $s_i$  represent the mean and standard deviation of examinee  $i$ 's current estimated proficiency distribution.

We propose to use Gauss-Hermite quadrature with a small number of quadrature points (approximately five). Experiments show that we can complete this calculation for 1,000 items in fewer than 5 milliseconds, making it computationally reasonable.

As with the binary items, we propose to ignore the slope parameters to even exposure and avoid a bias toward the items with better measurement.

### 1.3.4 Item Group Information Value

Item groups differ from individual items in that a set of items will be selected for administration. Therefore, the goal is to maximize information across the working theta distribution. As with the polytomous items, we propose to use Gauss-Hermite quadrature to estimate the expected information of the item group.

In the case of multiple-item groups

$$E[I_g(\theta)|m_i, s_i] = \frac{1}{J_g} \int \sum_{j=1}^{J_g} I_{g(j)}(t) \phi(t; m_i, s_i) dt$$

Where  $I_g(.)$  is the information from item group  $g$ ,  $I_{g(j)}$  is the information associated with item  $j \in g$ , for the  $J_g$  items in set  $g$ . In the case of polytomous items, we use the expected information, as described above.

## 2. ENTRY AND INITIALIZATION

At startup, the system will

- create a custom item pool;
- initialize theta estimates for the overall score and each score point; and
- insert embedded field-test items.

### 2.1 Item Pool

At test startup the system will generate a *custom item pool*, a string of item IDs for which the student is eligible. This item pool will include all items that

- are active in the system at test startup; and
- are not flagged as “access limited” for attributes associated with this student.

The list will be stored in ascending order of ID.

### 2.2 Adjust Segment Length

Custom item pools run the risk of being unable to meet segment blueprint minimums. To address this special case, the algorithm will adjust the blueprint to be consistent with the custom item pool. This capability becomes necessary when an accommodated item pool systematically excludes some content.

Let

$S$  be the set of top-level content constraints in the hierarchical set of constraints, each consisting of the tuple  $(name, min, max, n)$ ;

$C$  be the custom item pool, each element consisting of a set of content constraints  $B$ ;  
 $f, p$  integers represent item shortfall and pool count, respectively; and  
 $t$  be the minimum required items on the segment.

For each  $s$  in  $S$ , compute  $n$  as the sum of active operational items in  $C$  classified on the constraint.

$f = \text{summation over } S (min - n)$

$p = \text{summation over } S (n)$

if  $t - f < p$ , then  $t = t - f$

## 2.3 Initialization of Starting Theta Estimates

The user will supply five pieces of information in the test configuration:

1. A default starting value if no other information is available
2. An indication whether prior scores on the same test should be used, if available
3. Optionally, the test ID of another test that can supply a starting value, along with
4. Slope and intercept parameters to adjust the scale of the value to transform it to the scale of the target test
5. A constant prior variance for use in calculation of working EAP scores

## 2.4 Insertion of Embedded Field-Test Items

Each blueprint will specify

- the number of field-test items to be administered on each test;
- the first item position into which a field-test item may be inserted; and
- the last item position into which a field-test item may be inserted.

Upon startup, select randomly from among the field-test items or item sets until the system has selected the specified number of field-test items. If the items are in sets, the sets will be administered as a complete set, and this may lead to more than the specified number of items administered.

The probability of selection will be given by  $p_j = \frac{\sum_{j=1}^K K_j}{\sum_{j=1}^K a_j K_j} a_j K_j \frac{m}{N_j}$ , where

$p_j$  represents the probability of selecting the item;

$m$  is the targeted number of field-test items;

$N_j$  is the total number of active items in the field-test pool;

$K_j$  is the number of items in item set  $j$ ; and

$a_j$  is a user-supplied weight associated with each item (or item set) to adjust the relative probability of selection.

The  $a_j$  variables are included to allow for operational cases in which some items must complete field-testing sooner, or enter field-testing later. While using this parameter presents some statistical risk, not doing so poses operational risks.

For each item set, generate a uniform random number  $r_j$  on the interval  $\{0,1\}$ . Sort the items in ascending order by  $\frac{r_j}{p_j}$ . Sequentially select items, summing the number of items in the set. Stop the selection of field-test items once  $FTNMin \leq m \leq FTNMax = \sum_{j=0} K_j$ .

Next, each item is assigned to a position on the test. To do so, select a starting position within  $f - FTMax - FTMin$  positions from  $FTMin$ , where  $FTMax$  is the maximum allowable position for field-test items and  $FTMin$  is the minimum allowable position for field-test items.  $FTNMin$  and  $FTNMax$  refer to the minimum and maximum number of field-test items, respectively. Distribute the items evenly within these positions.

### 3. ITEM SELECTION

Exhibit 3 summarizes the item selection process. If the item position has been designated for a field-test item, administer that item. Otherwise, the adaptive algorithm kicks in.

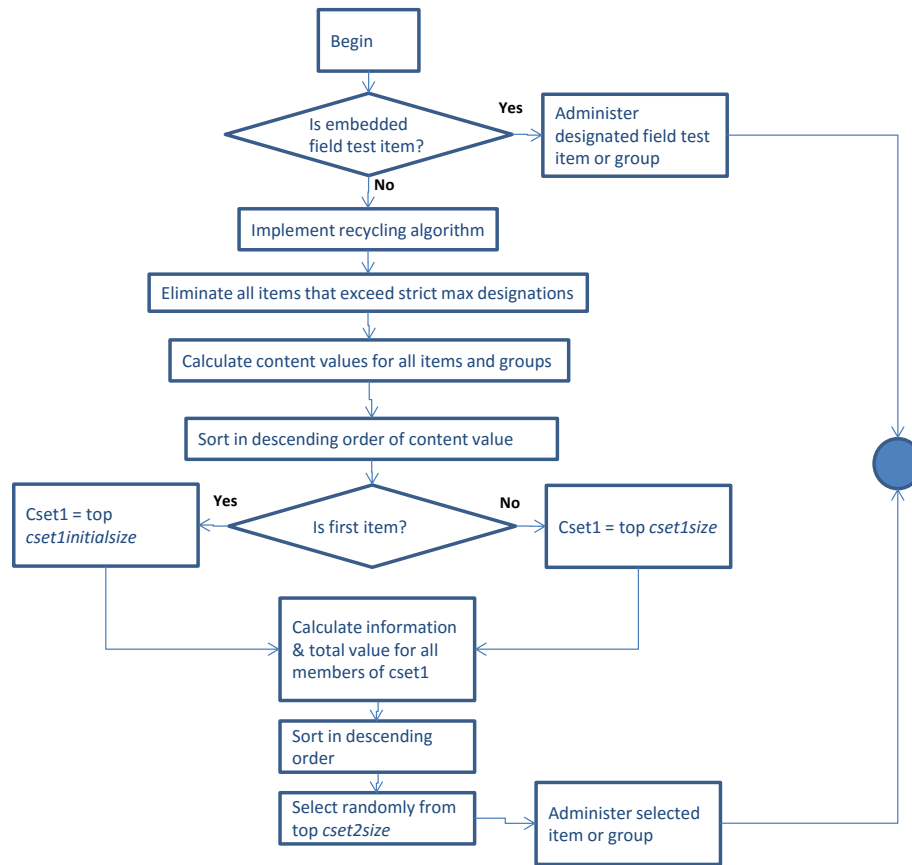


Exhibit 3: Summary of Item Selection Process

This approach is a “content first” approach designed to optimize match to blueprint. An alternative, “information first” approach, is possible. Under an information first approach, all items within a specified information range would be selected as the first set of candidates, and subsequent selection within that set would be based, in part, on content considerations. The engine is being designed so that future development could build such an algorithm using many of the calculations already available.

### 3.1 Trimming the Custom Item Pool

At each item selection, the active item pool is modified in four steps:

1. The custom item pool is intersected with the active item pool, resulting in a custom active item pool.
2. Items already administered on this test are removed from the custom active item pool.
3. Items that have been administered on prior tests are tentatively removed (see Section 3.2 below).

4. Items that measure content that has already exceeded a strict maximum are tentatively removed from the pool, removing entire sets containing items that meet this criterion.

## 3.2 Recycling Algorithm

When students are offered multiple opportunities to test, or when prior tests have been started and invalidated, students will have seen some of the items in the pool. The trimming of the item pool eliminates these items from the pool. It is possible that in such situations, the pool may no longer contain enough items to meet the blueprint.

Hence, items that have been seen on previous administrations may be returned to the pool. If there are not enough items remaining in the pool, the algorithm will recycle items (or item groups) with the required characteristic that is found in insufficient numbers. Working from the least recently administered group, items (or item groups) are reintroduced into the pool until the number of items with the required characteristics meets the minimum requirement. When item groups are recycled, the entire group is recycled rather than an individual item. Items administered on the current test are never recycled.

## 3.3 Adaptive Item Selection

Selection of items will follow a common logic, whether the selection is for a single item or an item group. Item selection will proceed in the following three steps:

1. Select Candidate Set 1 (*cset1*).
  - a. Calculate the content value of each item or item group.
  - b. Sort the item groups in descending order of content value.
  - c. Select the top *cset1size*, a user-supplied value that may vary by test.
2. Select Candidate Set 2 (*cset2*).
  - a. Calculate the information values for each item group in *cset1*.
  - b. Calculate the overall value of each item group in *cset1* as defined in Equation 1.
  - c. Sort *cset2* in descending order of value.
  - d. Select the top *cset2size* item groups, where *cset2size* is a user-supplied value that may vary by test.
3. Select the item or item group to be administered.
  - a. Select randomly from *cset2* with uniform probability.

Note that a “pure adaptive” test, without regard to content constraints, can be achieved by setting *cset1size* to the size of the item pool and  $w_2$ , the weight associated meeting content constraints in Equation 1, to zero. Similarly, linear-on-the-fly tests can be constructed by setting  $w_0$  and  $w_1$  to zero.

### 3.4 Selection of the Initial Item

Selection of the initial item can affect item exposure. At the start of the test, all tests have no content already administered, so the items and item groups have the same content value for all examinees. In general, it is a good idea to spread the initial item selection over a wider range of content values. Therefore, we define an additional user-settable value, *cset1initialsize*, which is the size of Candidate Set 1 on the first item only. Similarly, we define *cset2initialize*.

### 3.5 Exposure Control

This algorithm uses randomization to control exposure and offers several parameters that can be adjusted to control the tradeoff between optimal item allocation and exposure control. The primary mechanism for controlling exposure is the random selection from *CSET2*, the set of items or item groups that best meet the content and information criteria. These represent the “top *k*” items, where *k* can be set. Larger values of *k* provide more exposure control at the expense of optional selection.

In addition to this mechanism, we avoid a bias toward items with higher measurement precision by treating all items as though they measured with equal precision by ignoring variation in the slope parameter. This has the effect of randomizing over items with differing slope parameters. Without this step, it would be necessary to have other *post hoc* explicit controls to avoid the overexposure of items with higher slope parameters, an approach that could lead to different test characteristics over the course of the testing window.

## 4. TERMINATION

The algorithm will have configurable termination conditions. These may include

- administering a minimum number of items in each reporting category and overall;
- achieving a target level of precision on the overall test score;
- achieving a target level of precision on all reporting categories; and
- achieving a score insufficiently distant from a specified score with sufficient precision (e.g., less than two standard errors below proficient).
- We will define four user-defined flags indicating whether each of these is to be considered in the termination conditions (*TermCount*, *TermOverall*, *TermReporting*, *TermTooClose*). A fifth user-supplied value will indicate whether these are taken in conjunction or if satisfaction of any one of them will suffice (*TermAnd*). Reaching the minimum number of items is always a necessary condition for termination.

In addition, two conditions will each individually and independently cause termination of the test:

1. Administering the maximum number of items specified in the blueprint
2. Having no items in the pool left to administer



## A1. DEFINITIONS OF USER-SETTABLE PARAMETERS

This appendix summarizes the user-settable parameters in the adaptive algorithm.

<b>Parameter Name</b>	<b>Description</b>	<b>Entity Referred to by Subscript Index</b>
$w_0$	Priority weight associated with match to blueprint	N/A
$w_1$	Priority weight associated with reporting category information	N/A
$w_2$	Priority weight associated with overall information	N/A
$q_k$	Priority weight associated with a specific reporting category	reporting categories
$p_r$	Priority weight associated with a feature specified in the blueprint (These inputs appear as a component of the blueprint.)	features specified in the blueprint
$a$	Parameter of the function $h(.)$ that controls the overall information weight when the information target has not yet been hit	N/A
$b$	Parameter of the function $h(.)$ that controls the overall information weight after the information target has been hit	N/A
$c_k$	Parameter of the function $h(.)$ that controls the information weight when the information target has not yet been hit for reporting category $k$	reporting categories
$d_k$	Parameter of the function $h(.)$ that controls the information weight after the information target has been hit for reporting category $k$	reporting categories
cset1size	Size of candidate pool based on contribution to blueprint match	N/A
cset1initialsize	Size of candidate pool based on contribution to blueprint match for the first item or item set selected	N/A
cset2size	Size of final candidate pool from which to select randomly	N/A
cset2initialsize	Size of candidate pool based on contribution to blueprint match and information for the first item or item set selected	
$t_0$	Target information for the overall test	N/A
$t_k$	Target information for reporting categories	reporting categories
startTheta	A default starting value if no other information is available	N/A
startPrevious	An indication of whether previous scores on the same test should be used, if available	N/A
startOther	The test ID of another test that can supply a starting value, along with startOtherSlope	N/A
startOtherSlope	Slope parameter to adjust the scale of the value to transform it to the scale of the target test	N/A
startOtherInt	Intercept parameter to adjust the scale of the value to transform it to the scale of the target test	N/A
FTMin	Minimum position in which field-test items are allowed	N/A
FTMax	Maximum position in which field-test items are allowed	N/A
FTNMin	Target minimum number of field-test items	N/A
FTNMax	Target maximum number of field-test items	N/A
$a_j$	Weight adjustment for individual embedded field-test items used to increase or decrease their probability of selection	field-test items

<b>Parameter Name</b>	<b>Description</b>	<b>Entity Referred to by Subscript Index</b>
AdaptiveCut	The overall score cutscore, usually proficiency, used in consideration of <i>TermTooClose</i>	
TooCloseSEs	The number of standard errors below which the difference is considered “too close” to the adaptive cut to proceed. In general, this will signal proceeding to a final segment that contains off-grade items. Ugh.	
TermOverall	Flag indicating whether to use the overall information target as a termination criterion	N/A
TermReporting	Flag to indicate whether to use reporting category information target as a termination criterion	N/A
TermCount	Flag to indicate whether to use minimum test size as a termination condition	N/A
TermTooClose	Terminate if you are not sufficiently distant from the specified adaptive cut	
TermAnd	Flag to indicate whether the other termination conditions are to be taken separately or conjunctively	N/A

## A2. API

*This information is forthcoming.*

## A3. SUPPORTING DATA STRUCTURES

### **AIR Cautions and Caveats**

- Use of standard error termination conditions will likely cause inconsistencies between the blueprint content specifications and the information criteria will cause unpredictable results, likely leading to failures to meet blueprint requirements.
- The field-test positioning algorithm outlined here is very simple and will lead to deterministic placement of field-test items.